Relation Extraction

Justin T. Chiu

July 8, 2019

Abstract

Many recent relation extraction systems predict the relationship between an entity and value given the positions of their mentions in the text. This requires requires words to be annotated as mentions. Human annotation at the word level does not scale as the size of the text and the number of labels increases, as annotators must read every word. Automatic annotation methods allow the annotation process to scale, but may introduce noise due to incorrect annotations. In order to train a probabilistic information extraction model without mention annotations, we specify a model that, for each word, either chooses a triple from a knowledge base to explain or chooses to explain nothing.

1 Problem Statement

NOTE on RE vs KBP: KBP focuses on modeling the full database, whereas in RE the emphasis is on explaining extractions. KBP usually has a more complicated model over the KB, while RE sometimes makes simplifying assumptions and could make a poor KBP system. Example: if triple double appears in text and someone scored ¡ 10 PTS, then it's more likely they more likely had ¿ 10 REB. The difference between RE and KBP is subtle. Both fall under information extraction, and the models may overlap in many ways. In fact, latent variable model approaches to RE are usually able to perform KBP as well. However, the focus is usually not on having good aggregate extraction metrics, but rather about identifying where facts are in text. This leads to simplifying assumptions in the model of the KB itself. On the other hand, KBP moves the difficulty from explaining extractions to faithfully modeling the KB. Typically this results in a challenging inference task, as the structure of the KB may be very complex.

In relation extraction the goal is to extract facts from a passage of text. Systems must convert facts expressed in natural language into a form amenable to computation. Facts consist of three components: entities, relation types, and values. The challenge is to not only extract facts from text, but also justify the extractions by determining where those facts are mentioned.

A mention is a surface realization of an abstract object in text. In relation extraction we justify extractions by identifying fact mentions. As text is noisy, the realization of a fact may

be difficult to locate. We focus on locating fact mentions at the word level by identifying individual words as value mentions, rather than entity or type mentions.

The problem description is as follows. We focus on the domain of basketball summaries: given a written summary of a basketball game $x = x_1, \ldots, x_I$ we model the aligned box score $\{(e_j, t_j, v_j)\}_{j=1}^J$ consisting of entities e_j, relation types t_j, and all values v_j. The set of facts is our knowledge base (KB). Let $d = \{(e_j, t_j)\}_{j=1}^J$ and $v = \{v_j\}_{j=0}^J$. The KB (d, v) can be viewed as a data table where d defines a flattened representation of the rows and columns and v gives the values of the cells. For example, we may have $d = \{(e_1 = \text{John}, t_1 = \text{Points}), (e_2 = \text{John}, t_2 = \text{Rebounds}), \ldots\}$ with $v = \{v_1 = 19, v_2 = 12, \ldots\}$ aligned to the brief summary $x = \text{John scored 19 points}$. Our goal is to locate and extract facts from x.

(move somewhere else) We are primarily concerned with the scenario where we have an overcomplete KB schema with respect to a specific passage of text. This fits many scenarios in real world applications: we may have many entity and type pairs in our data table, but a summary may discuss only a small salient subset of players and statistics.

As our goal is to locate fact mentions, modeling just the KB (d, v) given the text x is not sufficient. Suppose we had a model that directly filled in $v \mid x, d$. Without further assumptions, it would be difficult to separate the contribution of the text from the marginal distribution of values. If we are indeed interested in the extraction of facts from text, then we must be able to answer whether facts are supported in text.

We propose a model that first identifies value mentions at the word level, aligns those mentions to an entity and relation type in order to obtain a fact, then aggregates word level decisions to resolve conflicts.

(shoot, this is incorrect. we actually care about granular extraction) Since we only observe a text and associated KB, we must We reduce the construction of Δ to predicting, for every $j \in 1, \ldots, J$, the value v_j corresponding to row and column of our KB given by s_j. We cast relation extraction as KB completion, where we fill in the cells of a data table. In the following section we propose a model for the distribution over values given text and a value-less KB $p(v \mid x, s)$.

2   Model

We define a graphical model that performs extraction with justification. The model first extracts information at the word level, then aggregates its choices for each word into an extraction at the sequence level.

The word level extraction process has three steps. For each index $i \in 1, \ldots, I$ we perform

1. Value mention identification: Given a sequence of words x, we identify whether each word is a value mention with $p(m \mid x) = \prod_i p(m_i \mid x)$. Each $m_i \in \{0, 1\}$. Not every word in a mention must be identified; it suffices to find at least one word in a value mention.

2. Alignment: Each value mention is then aligned to a record in the knowledge base with $p(a \mid x, d) = \prod_i p(a_i \mid x, d)$, We align the word x_i by choosing who (the entity) and what (the relation type) generate the possible value mention at index i. In particular,

$a_i = j$ denotes the alignment to the record r_j with $a_i \in 1, \ldots, J$. We assume that each value mention aligns to a single record.

3. Translation: All value mentions are translated into a value from the KB schema with $p(z \mid x) = \prod_i p(z_i \mid x)$, with $z_i \in \mathcal{V}$.

Finally, we aggregate the word level information at the sequence level in order to give a single distribution over the record values for x.

4. Aggregation $p(v \mid z, a, m, d) = \prod_j p(v_j \mid z, a, m, d_j)$: Given the word level values z, alignments a, value mention decisions m, and data table d we choose the sequence level value v_j.


Figure 1: Our model predicts word-level values and alignments then aggregates those choices over all indices i to predict values at the KB level.

This gives us the following factorization of the relation extraction system:

$$
\begin{aligned}
p(v \mid x, d) &= \sum_{z,a,m} p(v, z, a, m \mid x, d) \\
&= \sum_{z,a,m} p(v \mid z, a, m, x, d) p(z, a, m \mid x, d) \\
&= \sum_{z,a,m} \left(\prod_j p(v_j \mid z, a, m, x, d) \right) \left(\prod_i p(z_i \mid x) p(a_i \mid x, d) p(m_i \mid x) \right)
\end{aligned}
\tag{1}
$$

2.1 Parameterization

Our model has four steps: mention identification, mention alignment, mention translation, and aggregation. We parameterize the conditional distributions of each step below.

Let $\mathbf{h}_i \in \mathbb{R}^d$ be a contextual embedding of the word x_i, and E an embedding function that maps entities and types to vectors in $\mathbb{R}^{d'}$.

1. Identification: We use the contextual embedding to directly predict whether a word is part of a value mention.

$$
p(m_i \mid x) \propto \exp(W_m \mathbf{h}_i), W_m \in \mathbb{R}^{2 \times d}
$$

2. Alignment: We decompose the alignment distribution into a distribution over entities $p(\epsilon_i \mid x, d)$ and types $p(\tau_i \mid x, d)$.

$$p(a_i \mid x, d) = p(\epsilon_i \mid x, d)p(\tau_i \mid x, d)$$
$$p(\epsilon_i \mid x, d) \propto \exp(E(e_{\epsilon_i})^T W_e \mathbf{h}_i)$$
$$p(\tau_i \mid x, d) \propto \exp(E(\tau_{a_i})^T W_t \mathbf{h}_i)$$

with $W_e \in \mathbb{R}^{d' \times d}, W_t \in \mathbb{R}^{d' \times d}$.

3. Translation: We use the contextual embedding to translate a word into a value.

$$p(z_i \mid x) \propto \exp(W_z \mathbf{h}_i), W_z \in \mathbb{R}^{|\mathcal{V}| \times d}$$

4. Aggregation: If there exists an index that is a mention and is also aligned to r_j we allow it to vote on the value v_j, otherwise we ignore the text and use a prior distribution over values $p(v_j \mid d_j) \propto \exp(E(v_j)^T W_v[E(d_j)])$.

$$p(v_j \mid z, a, m, d) \propto \begin{cases} \prod \exp(\psi(v_j, z_i, a_i, m_i, d)), & \exists i, m_i = 1 \wedge a_i = j \\ \exp(E(v_j)^T W_v[E(d_j)]), & \text{otherwise} \end{cases}$$
$$\psi(v_j, z_i, a_i, m_i, d) = 1(v_j = z_i, a_i = j, m_i = 1)$$

3 Training and Inference

To train a latent variable model, we must marginalize over the unobserved RVs and maximize the likelihood of the observed. Ideally, we would optimize the following objective

$$\log p(v \mid x, d) = \log \sum_{z, a, m} p(v, z, a, m \mid x, d) \tag{2}$$

However, maximizing $\log p(v \mid x, d)$ directly is very expensive for this model as the summation over z, a, m is intractable. The summation over z, a, m has computational complexity $O((|\mathcal{V}| \cdot J \cdot 2)^I)$, which is exponential in the length of the text. Additionally, the size of the KB J may be large as well.

We therefore resort to approximate inference, specifically amortized variational inference.

3.1 Inference network

Our first approach is to introduce an inference network $q(z, a, m \mid v, x, d)$ and optimize the following lower bound on the marginal likelihood with respect to the parameters of both p and q:

$$\log p(v \mid x) \geq \mathbb{E}_{q(z, a, m \mid v, x, d)} \left[\log \frac{p(v, z, a, m \mid x, d)}{q(z, a, m \mid v, x, d)} \right] \tag{3}$$

We propose to parameterize $q(z, a, m \mid v, x, d)$ as follows. We decompose

$$\begin{aligned} q(z, a, m \mid v, x, d) &= q(z \mid a, v, x)q(a \mid v, x, d)q(m \mid v, x) \\ &= \prod_i q(z_i \mid a, v, x)q(a_i \mid v, x, d)q(m_i \mid v, x) \end{aligned} \tag{4}$$

The conditional distributions of our inference network are very similar to the relation extraction model, but they condition on the values v.

Let $\mathbf{h}_i \in \mathbb{R}^d$ be a contextual embedding of the word x_i. We use attention weights over records to get a weighted representation of the records of the KB for each index i:

$$\mathbf{g}_{r_j} = [E(d_j), E(v_j)]$$
$$\alpha_j \propto \exp(\mathbf{g}_{r_j}^T W_\alpha \mathbf{h}_i)$$

The inference network is given by

1. The value mention model $q(m_i \mid v, x)$ has access to the values v from the KB, which it conditions on when detecting value mentions.

$$p(m_i \mid v, x) = W'_m \text{MLP}([\sum_j \alpha_j \cdot \mathbf{g}_{r_j}, \mathbf{h}_i]), W'_m \in \mathbb{R}^{2 \times d}$$

2. The alignment model $q(a_i \mid v, x, d)$ uses a contextual representation of each x_i and chooses a record. In contrast to $p(a \mid x, d)$, this model has access to values as well. We use the attention weights to parameterize the distribution $p(a_i \mid x) = \alpha_{a_i}$.

3. The translation model $q(z_i \mid a, v, x) = 1(z_i = v_{a_i})$ conditions on the alignments a and ensures the chosen z is consistent with the alignments.

One concern is that the model may learn to never rely on the text for extraction, setting $m_i = 0$ at every index. We can avoid this by initializing $q(z)$ to ensure that for words $x \in \mathcal{V}$ we have $q(z = x)$ is high, biasing the translation model towards transliteration at the start of training.

3.2   Approximate the posterior of a generative model

Alternatively, we may introduce a generative model of text $q(x, v)$ that inverts the relation extraction model $p(v \mid x)$ and optimize the following objective:

$$\log q(x, v) - KL[p(v, z, a, m \mid x)||q(v, z, a, m \mid x)] \tag{5}$$

which entails training the generative model of text while approximating its posterior with the information extraction system.

decompose the training of our extraction system $p(v \mid x)$ into two stages: In the first stage we train $p(z, a, m \mid x, d)$ to approximate the posterior of a conditional model of text given a complete KB $q(x, z, a, m \mid v, d)$. This has the benefit of allowing us to exert control over where value mentions are detected through our design of the text model q.

In the second stage, we have two choices: a) train $p(v \mid z, a, m, x, d)$ to approximate the posterior of a full generative model of text and the values of KB $q(x, v \mid d)$. b) train $p(v \mid z, a, m, x, d)$ using the following lower bound:

$$\log p(v \mid x, d) \geq \mathbb{E}_{p(z, a, m | x, d)} [\log p(v \mid z, a, m, x, d)] \tag{6}$$

Ideally the bound in Eqn. 6 should not be looser than the one presented in Eqn. 3, as conditioning on the observed values of a KB should not reduce the entropy of a good alignment model.

How do we split the gradient over time?

4   Evaluation

Although we have a model over the values of all records, evaluation does not include the final distribution over all record values. As we assumed that the KB contained a superset of the facts contained in a sequence of text, we are evaluate whether the model can discover the subset of facts that are expressed in the text. We therefore perform extraction by using the marginal distributions $q(z), q(a), q(c)$ to value mentions as well as entities and types, giving us facts.

References