

Relation Extraction

Justin T. Chiu

July 11, 2019

Abstract

Many recent relation extraction systems predict the relationship between an entity and value given the positions of their mentions in the text. This requires words to be annotated as mentions. Human annotation at the word level does not scale as the size of the text and the number of labels increases, as annotators must read every word. Automatic annotation methods allow the annotation process to scale, but may introduce noise due to incorrect annotations. In order to train a probabilistic information extraction model without mention annotations, we specify a model that, for each word, either chooses a triple from a knowledge base to explain or chooses to explain nothing.

1 Problem Statement

In relation extraction the goal is to extract facts from a passage of text. Systems must convert facts expressed in natural language into a form amenable to computation. Facts consist of three components: entities, relation types, and values. The challenge is to not only extract facts from text, but also justify the extractions by determining where those facts are mentioned.

A mention is a surface realization of an abstract object in text. In relation extraction we justify extractions by identifying fact mentions. As text is noisy, the realization of a fact may be difficult to locate. We focus on locating fact mentions at the word level by identifying individual words as value mentions, rather than entity or type mentions.

The problem description is as follows. We focus on the domain of basketball summaries: given a written summary of a basketball game $x = x_1, \dots, x_I$ we model the aligned box score $\{(e_j, t_j, v_j)\}_{j=1}^J$ consisting of entities e_j , relation types t_j , and all values $v_j \in \mathcal{V}$. The set of facts is our knowledge base (KB). Let $d = \{(e_j, t_j)\}_{j=1}^J$ and $v = \{v_j\}_{j=0}^J$. The KB (d, v) can be viewed as a data table where d defines a flattened representation of the rows and columns and v gives the values of the cells. For example, we may have $d = \{(e_1 = \text{John}, t_1 = \text{Points}), (e_2 = \text{John}, t_2 = \text{Rebounds}), \dots\}$ with $v = \{v_1 = 19, v_2 = 12, \dots\}$ aligned to the brief summary $x = \text{John scored 19 points}$. Our goal is to locate and extract facts from x .

As our goal is to locate fact mentions, modeling just the KB (d, v) given the text x is not sufficient. We therefore propose a model that first identifies words as value mentions, aligns those mentions to an entity and relation type in order to obtain a fact, then aggregates word level decisions to resolve conflicts.

2 Model

We define a graphical model that performs extraction with justification. The model first extracts information at each word, then aggregates its choices for each word into an extraction at the sequence level.

The word level extraction process has three steps. For each index $i \in 1, \dots, I$ we perform

1. Value mention identification: Given a sequence of words x , we identify whether each word is a value mention with $p(m \mid x) = \prod_i p(m_i \mid x)$. Each $m_i \in \{0, 1\}$. Not every word in a mention must be identified; it suffices to find at least one word in a value mention.
2. Alignment: Each value mention is then aligned to a record in the knowledge base with $p(a \mid x, d) = \prod_i p(a_i \mid x, d)$. We align the word x_i by choosing who (the entity) and what (the relation type) generate the possible value mention at index i . In particular, $a_i = j$ denotes the alignment to the record r_j with $a_i \in 1, \dots, J$. We assume that each value mention aligns to a single record.
3. Translation: All value mentions are translated into a value from the KB schema with $p(z \mid x) = \prod_i p(z_i \mid x)$, with $z_i \in \mathcal{V}$.

Finally, we aggregate the word level information at the sequence level in order to give a single distribution over the record values for x .

4. Aggregation $p(v \mid z, a, m, d) = \prod_j p(v_j \mid z, a, m, d_j)$: Given the word level values z , alignments a , value mention decisions m , and data table d we choose the sequence level value v_j .

This gives us the following factorization of the relation extraction system:

$$\begin{aligned}
 p(v \mid x, d) &= \sum_{z, a, m} p(v, z, a, m \mid x, d) \\
 &= \sum_{z, a, m} p(v \mid z, a, m, x, d) p(z, a, m \mid x, d) \\
 &= \sum_{z, a, m} \left(\prod_j p(v_j \mid z, a, m, x, d) \right) \left(\prod_i p(z_i \mid x) p(a_i \mid x, d) p(m_i \mid x) \right)
 \end{aligned} \tag{1}$$

2.1 Parameterization

Our model has four steps: mention identification, mention alignment, mention translation, and aggregation. We parameterize the conditional distributions of each step below.

Let $\mathbf{h}_i \in \mathbb{R}^d$ be a contextual embedding of the word x_i , and E an embedding function that maps entities and types to vectors in $\mathbb{R}^{d'}$.

1. Identification: We use the contextual embedding to directly predict whether a word is part of a value mention.

$$p(m_i \mid x) \propto \exp(W_m \mathbf{h}_i), W_m \in \mathbb{R}^{2 \times d}$$




Figure 1: Our model predicts word-level values and alignments then aggregates those choices over all indices i to make a single decision for each value. Each word has the following latent variables: the mention $m_i \in \{0, 1\}$ indicates whether word x_i is a value mention, the alignment a_i gives what fact x_i aligns to, and the value z_i gives the canonical value that x_i translates to.

2. Alignment: We decompose the alignment distribution into a distribution over entities $p(\epsilon_i | x, d)$ and types $p(\tau_i | x, d)$.

$$\begin{aligned} p(a_i | x, d) &= p(\epsilon_i | x, d)p(\tau_i | x, d) \\ p(\epsilon_i | x, d) &\propto \exp(E(e_{\epsilon_i})^T W_e \mathbf{h}_i) \\ p(\tau_i | x, d) &\propto \exp(E(\tau_{a_i})^T W_t \mathbf{h}_i) \end{aligned}$$

with $W_e \in \mathbb{R}^{d' \times d}$, $W_t \in \mathbb{R}^{d' \times d}$.

3. Translation: We use the contextual embedding to translate a word into a value.

$$p(z_i | x) \propto \exp(W_z \mathbf{h}_i), W_z \in \mathbb{R}^{|\mathcal{V}| \times d}$$

4. Aggregation: If there exists an index that is a mention and is also aligned to r_j we allow it to vote on the value v_j , otherwise we ignore the text and use a prior distribution over values $p(v_j | d_j) \propto \exp(E(v_j)^T W_v [E(d_j)])$.

$$\begin{aligned} p(v_j | z, a, m, d) &\propto \begin{cases} \prod \exp(\psi(v_j, z_i, a_i, m_i, d)), & \exists i, m_i = 1 \wedge a_i = j \\ \exp(E(v_j)^T W_v [E(d_j)]), & \text{otherwise} \end{cases} \\ \psi(v_j, z_i, a_i, m_i, d) &= 1(v_j = z_i, a_i = j, m_i = 1) \end{aligned}$$

3 Training and Inference

To train a latent variable model, we must marginalize over the unobserved RVs and maximize the likelihood of the observed. Ideally, we would optimize the following objective

$$\log p(v | x, d) = \log \sum_{z, a, m} p(v, z, a, m | x, d) \quad (2)$$

However, maximizing $\log p(v \mid x, d)$ directly is very expensive for this model as the summation over z, a, m is intractable. The summation over z, a, m has computational complexity $O((|\mathcal{V}| \cdot J \cdot 2)^I)$, which is exponential in the length of the text. Additionally, the size of the KB J may be large as well.

We therefore resort to approximate inference, specifically amortized variational inference.

3.1 Inference network

Our first approach is to introduce an inference network $q(z, a, m \mid v, x, d)$ and optimize the following lower bound on the marginal likelihood with respect to the parameters of both p and q :

$$\log p(v \mid x) \geq \mathbb{E}_{q(z, a, m \mid v, x, d)} \left[\log \frac{p(v, z, a, m \mid x, d)}{q(z, a, m \mid v, x, d)} \right] \quad (3)$$

We propose to parameterize $q(z, a, m \mid v, x, d)$ as follows. We decompose

$$\begin{aligned} q(z, a, m \mid v, x, d) &= q(z \mid a, v, x) q(a \mid v, x, d) q(m \mid v, x) \\ &= \prod_i q(z_i \mid a, v, x) q(a_i \mid v, x, d) q(m_i \mid v, x) \end{aligned} \quad (4)$$

The conditional distributions of our inference network are very similar to the relation extraction model, but they condition on the values v .

Let $\mathbf{h}_i \in \mathbb{R}^d$ be a contextual embedding of the word x_i . We use attention weights over records to get a weighted representation of the records of the KB for each index i :

$$\begin{aligned} \mathbf{g}_{r_j} &= [E(d_j), E(v_j)] \\ \alpha_j &\propto \exp(\mathbf{g}_{r_j}^T W_\alpha \mathbf{h}_i) \end{aligned}$$

The inference network is given by

1. The value mention model $q(m_i \mid v, x)$ has access to the values v from the KB, which it conditions on when detecting value mentions.

$$p(m_i \mid v, x) = W'_m \text{MLP}([\sum_j \alpha_j \cdot \mathbf{g}_{r_j}, \mathbf{h}_i]), W'_m \in \mathbb{R}^{2 \times d}$$

2. The alignment model $q(a_i \mid v, x, d)$ uses the attention weights to parameterize the alignment $p(a_i \mid x) = \alpha_{a_i}$.
3. The translation model $q(z_i \mid a, v, x) = 1(z_i = v_{a_i})$ conditions on the alignments a and ensures the chosen z is consistent with the alignments.

We use the score function gradient estimator to perform gradient ascent on the objective in Equation 4.

One concern is that the model may learn to never rely on the text for extraction, setting $m_i = 0$ at every index. We can avoid this by initializing $q(z)$ to ensure that for words $x \in \mathcal{V}$ we have $q(z = x)$ is high, biasing the translation model towards transliteration at the start of training.

4 Evaluation

As we assumed that the KB contained a superset of the facts contained in a sequence of text, we are interested in evaluating whether the model can discover and locate the subset of facts that are expressed in the text.

We evaluate first by outputting a

We therefore perform extraction by using the marginal distributions $q(z), q(a), q(c)$ to value mentions as well as entities and types, giving us facts.

References